

# Agentic AI - Complete Study Notes

## 1. Transformer Architecture

The Transformer is the foundational architecture behind modern LLMs like GPT, Claude, LLaMA, and Qwen. It processes input tokens in parallel instead of sequentially, using \*\*self-attention\*\*.

Architecture flow: Input Tokens → Embedding → Positional Encoding → Encoder/Decoder Layers → Attention → Output Tokens. Each layer has Multi-Head Attention, Add & Norm, and Feed-Forward layers.

## 2. Attention, Q, K, V Vectors, SoftMax

Each token produces Query (Q), Key (K), and Value (V) vectors using learned matrices. Attention =  $\text{Softmax}((Q \times K^T) / \sqrt{dk}) \times V$ . Softmax normalizes attention scores into probabilities that highlight important tokens.

## 3. Pre-training Objective - Next Token Prediction

LLMs are pre-trained using next-token prediction: given tokens  $[t_1, t_2, \dots, t_n]$ , predict  $t_{n+1}$ . It builds contextual understanding through billions of token predictions.

## 4. Decoding Methods - Temperature, Top P, Top K, Greedy

- Temperature: controls randomness (low = deterministic, high = creative).
- Top-K: samples from top K most probable tokens.
- Top-P (nucleus): samples tokens whose cumulative probability  $\leq P$ .
- Greedy Decoding: always selects token with highest probability.

## 5. Tokenization

Text → Subword Units → IDs (tokens). Tokenizer ensures consistent mapping between text and model input. Example: “playing” → ['play', ‘ing’].

## 6. Supervised Finetuning & Instruction Tuning

Fine-tuning aligns LLMs with human instructions using supervised datasets (question → answer). It helps transform base LLM into instruction-following assistants like ChatGPT.

## 7. Prompting Techniques

- Zero-shot: no examples, only prompt.
- Few-shot: includes few examples for context.
- Chain-of-Thought (CoT): guides model step-by-step reasoning.
- Prompt Templates: System Prompt (sets behavior), User Prompt (task input).

## 8. RAG (Retrieval-Augmented Generation)

RAG enhances LLMs by combining external knowledge. Steps: Parsing → Chunking (Overlap/Semantic) → Embedding → Store in Vector DB → Query Embedding → Retrieval → Reranking → LLM Response Generation.

## 9. LangChain

LangChain is a Python framework that connects LLMs with tools, memory, and data sources. It simplifies building conversational agents, RAG pipelines, and chains of logic.

## 10. Gradio / Streamlit

Frameworks for creating user interfaces for AI models. • Gradio: Quick web demos with few lines of Python. • Streamlit: Interactive dashboards for ML apps.

## 11. FastAPI & Pydantic

FastAPI is a high-performance web framework for serving ML/LLM apps. Pydantic validates input/output using Python models. Used for API deployment, logging, and error handling.

## 12. LLaMA, DeepSeek, Claude, Qwen

• LLaMA (Meta): Open-source transformer models optimized for efficiency. • DeepSeek: Chinese LLM focused on reasoning and coding. • Claude (Anthropic): Family - Opus, Sonnet, Haiku (different model sizes for reasoning vs speed). • Qwen (Alibaba): Multilingual and code-capable transformer model family.

## 13. Agents & Tools

Agent = LLM + Tools (e.g., Web search, Python exec, APIs). They can act autonomously, plan tasks, and use memory.

## 14. Multi-Agent Architectures

Multiple agents collaborate. • Supervisor Agent: assigns tasks. • Worker/Swarm Agents: perform subtasks. • Handoffs: communication between agents. State management ensures context sharing.

## 15. Memory Types

• Episodic Memory: Stores recent conversations. • Semantic Memory: Stores structured knowledge. • Long-term memory ensures context continuity across sessions.

## 16. LangGraph

LangGraph provides visual graphs for defining agent workflows. Useful for building complex, stateful agent interactions.

## **17. MCP (Model Context Protocol)**

MCP is a framework for connecting AI models and external tools (APIs, databases). Building an MCP server: define tool schema → register endpoints → connect to model (e.g., Claude Desktop).

## **18. Developer Tools**

- Cursor: AI code editor for LLM pair programming.
- Claude Desktop: Anthropic's local chat & integration tool.
- GitHub Copilot: AI code assistant for IDEs.
- Lovable: UI-based app builder for AI workflows.

## **19. Memories**

Persistent storage of user context across sessions. Used to personalize LLM interactions and improve long-term understanding.